

popsim: Reproducing Elements of Billion-Agent Social Simulators on a Single Machine, with a Behavioral-Adherence Benchmark for Persona Agents on Sentiment140

Micah Stubbs
San Francisco, CA
micah@casemirror.ai

Yvonne Chen
San Francisco, CA
yc.yvonnechen@gmail.com

August 2026 — draft v0.2

Abstract

Population-scale social simulators such as Light Society (one billion agents) and MatrAIX (8.3 billion personas) require infrastructure most researchers cannot access. We show that their core methodology — human-grounded personas, tiered compute, and ground-truth validation — runs on one commodity machine, using Sentiment140: 1.6M public tweets in under 300 MB of memory. We mine personas for the 6,245 users with at least 20 tweets each, and a mixture-of-models ladder cuts simulation cost $702.7\times$ versus an all-LLM population. Validating on the real @-mention graph produces a negative result: DeGroot opinion diffusion ($r = 0.54$) underperforms per-user persistence ($r = 0.73$) at predicting future sentiment — disposition beats diffusion. Sustained one-way (parasocial) attachments prove as numerous as strong mutual ties, motivating a near-zero-cost broadcast tier. We also release a fixed behavioral-adherence benchmark: per-user temporal 80/20 splits scored on 42,998 held-out tweets, where a zero-parameter persona baseline reaches 63.1% and MatrAIX’s reported 91.5%, under a differing protocol, frames the gap. Code, artifacts, and an interactive explainer are public.

1 Introduction

Two recent systems define the frontier of population-scale social simulation. Light Society [1] advances “earth-scale” simulation: over 10^9 agents whose states evolve through LLM-powered simulation operations, made tractable by a mixture-of-models engine combining full LLMs with distilled surrogates, and validated through trust-game and opinion-diffusion experiments. MatrAIX [2] contributes population realism of a different kind: 8.3 billion persona records spanning 1,290

attribute dimensions, with a curated million-persona subset grounded partly in real human profiles, and a headline evaluation showing persona agents express or appropriately suppress declared behaviors in 91.5% of trials.

Both systems presume compute, data pipelines, and engineering budgets beyond most academic and hobbyist settings. This paper asks a deliberately modest question: *which elements of this research program can be reproduced on one machine, over one public dataset, in an evening?* Our answers are embodied in **popsim**, built at a public build night, with all analysis running in-memory on a single node.¹

Our contributions are:

1. **A single-machine, small-compute reproduction of elements of Light Society and MatrAIX:** human-grounded persona mining (§4); a mixture-of-models scale ladder with a measured $702.7\times$ cost reduction (§5); and validation of simulated dynamics against held-out ground truth on a real interaction graph (§6), including a negative result we believe is informative for simulator design.
2. **A new behavioral-adherence evaluation for persona agents on an existing public dataset** (§8): a fixed, reproducible protocol on Sentiment140 with a zero-parameter baseline of 63.1%, framed against MatrAIX’s (protocol-distinct) 91.5% bar.
3. **An empirical account of relationship strength and parasocial structure** in the dataset’s mention graph (§7), with a concrete

¹Interactive explainer: <https://popsim.micahstubbs.ai>.
Code and artifacts: <https://github.com/micahstubbs/pop-world-model-sim>.

08 - THE SUBSTRATE

Why Sentiment140

The event's two papers — *Light Society* (10⁹ agents) and *MatrAIx* (8.3B personas) — both stand on **human-grounded persona data**. Since the 2023 X API shutdown killed tweet-ID hydration, the practical question is: which full-text corpus gives the most distinct humans with enough history each to ground a persona — on one laptop?

CANDIDATE	USERS WITH ≥20 TWEETS	VERDICT
Sentiment140 (2009)	6,245 — measured	selected: 295 MB in RAM
Community Archive	~ 363 accounts total	full-corpus access paused
Archive Team 1% stream	~ 8	users rarely recur 20+
tweet-ID corpora (COVID etc.)	n/a	hydration is dead

All numbers on this page were computed locally from the downloaded corpus. Every tweet shown is real, public, and 17 years old — except the ones clearly marked synthetic.

Figure 1: Dataset selection in the interactive explainer (<https://popsim.micahstubbles.ai/#data>): candidate corpora compared on the number of distinct users with at least 20 tweets.

architectural implication: population simulators need a broadcast tier, and it is nearly free.

2 Related Work

Population-scale simulation. *Light Society* [1] formalizes society simulation as structured transitions of agent and environment states governed by LLM-powered operations, scaling to 10⁹ agents grounded in World Values Survey demographic profiles via a mixture-of-models engine. *MatrAIx* [2] builds an 8.3B-record persona database (599,847 human-grounded profiles among the curated subset) and evaluates behavioral adherence across four environments. Generative agent-based modeling more broadly descends from Park et al.’s generative agents [3], which demonstrated believable social behavior from memory-augmented LLM agents in a small sandbox town.

Data. Sentiment140 [4] provides 1.6M tweets from April–June 2009, distantly labeled for sentiment by emoticons. It remains one of the few large fully-public full-text tweet corpora: since the 2023 restriction of the X API, tweet-ID “hydration” workflows that most academic Twitter datasets rely on are effectively defunct, making full-text corpora the practical substrate for new work. TweetEval [5] standardizes tweet classification benchmarks and supplies pretrained classifiers usable as realism scorecards. Kwak et al. [6] established the large-scale structural view of Twitter; Horton and Wohl [7] introduced the parasocial-interaction construct we operationalize; DeGroot [8] supplies the consensus-formation model we test.




Figure 2: Per-user tweet counts are heavy-tailed. The cohort threshold (≥ 20 tweets) admits 6,245 users covering 226,202 tweets.

3 Dataset and Cohort

Sentiment140 loads from one CSV into 295 MB of memory: 1,600,000 tweets, 659,775 distinct users, 48 days, balanced 800K/800K negative/positive labels. Figure 2 shows the heavy-tailed per-user activity distribution. We define the **persona cohort** as the 6,245 users with at least 20 tweets (226,202 tweets; mean 36.2, median 28 per user; mean per-user activity span ~ 50 days). The corpus retains a strong diurnal rhythm (Figure 3) — peak 23:00, trough 13:00 PDT — usable for calibrating agent posting schedules. Selection among candidate corpora was itself empirical: the Community Archive holds only 363 accounts; the Archive Team 1% stream rarely observes a user 20 times; ID-only corpora cannot be hydrated. Sentiment140 maximizes the number of distinct humans with enough history to ground a persona while fitting in RAM.

4 Persona Mining

Following MatrAIx’s human-grounded persona construction — at hackathon scale — each cohort user becomes a persona card computed from their real timeline: sentiment disposition (share of positive tweets), verbosity (words/tweet), mention rate, peak posting hour, activity span, and characteristic vocabulary. Unlike survey-grounded profiles, these personas carry temporal behavior (posting cadence, diurnal phase) directly usable to drive simulation ticks. The cohort skews positive (59% vs. the corpus’s balanced 50%), a selection effect worth modeling rather than ignoring: highly active users differ from the population they are




Figure 3: Diurnal activity cycle across the corpus (all 1.6M tweets).

drawn from.

5 The Mixture-of-Models Scale Ladder

Light Society’s central efficiency device is heterogeneous agent cost: full LLMs only where fidelity demands them. We instantiate the same ladder sized to one node (Figure 4): **full LLM agents** for the 927 users with ≥ 50 tweets (enough signal to condition a frontier model); **distilled surrogates** (persona-conditioned small models) for the 5,318 users with 20–49 tweets; and a **statistical background** tier (rate and sentiment samplers) for the remaining 653,530 users. Under a relative cost model of 1000 : 1 : 0.01 units per agent-tick, the ladder costs 9.39×10^5 units per tick versus 6.60×10^8 for an all-LLM population — a 702.7 $\times$ reduction, with zero fidelity loss on the users who carry most of the observable signal.

6 Does Sentiment Diffuse? A Validated Negative Result

The cohort contains a real directed interaction graph: 31,515 @-mention events between cohort members. Light Society validates simulations with opinion-diffusion experiments; we run the corresponding test *against ground truth*. Each user is seeded with their first-half sentiment; DeGroot mixing [8] ($s_{t+1} = 0.7s_t + 0.3\bar{s}_N$, 10 steps, edge weights = mention counts) evolves states over the real graph; we then ask which predictor better recovers each user’s *second-half* sentiment (users with ≥ 5 tweets in each half; $n = 5,463$).

Persistence wins (Figure 5): $r = 0.73$ for a user’s own past versus $r = 0.54$ after diffusion; restricting to graph-connected users widens the gap (0.72 vs. 0.40). Neighbor-mixing strictly destroys predictive information at this horizon. **Disposition beats diffusion:** over 48 days, who a person is outweighs who they talk to. For simulator design the implication is direct — a population world model that over-weights social contagion will be *less* accurate than one grounded in stable personas, which is precisely why persona grounding (§4) matters.

7 Relationship Strength and Parasocial Structure

To model whom agents should talk to, we quantify relationship structure over the full mention graph (all 1.6M tweets). For unordered pairs where both users author tweets in the sample (so silence is observed, not censored; $n = 293,212$ pairs), we score reciprocity $r = \min(w_{ab}, w_{ba}) / \max(w_{ab}, w_{ba})$ and strength $\log(1+\text{total}) \cdot \log(1+\text{days}) \cdot (1+r)$, and classify pairs by mutuality and volume (Figure 6, left). Genuine relationships are rare: 77% of pairs are one-off contacts. Strikingly, sustained one-way attachments (*unidirectional strong*: ≥ 5 mentions, zero returned; 1,957 pairs) are as numerous as strong mutual bonds (1,901 pairs) — for every observable friendship there is roughly one unrequited attachment.

At the account level, among in-sample targets with ≥ 20 distinct mentioners, 408 reciprocate under 5% of their audience — operationally *parasocial* [7] — including 2009-era celebrities (*wossy*, *ashleytisdale*, *tomfelton*) who tweet but do not talk back, while 163 *community hubs* of comparable audience reciprocate 30–53%. A further 524 broadcast targets (e.g. *jonasbrothers*: 1,901 distinct fans in 48 days) never author a sampled tweet, so their reciprocity is unobservable and we label rather than classify them. By volume, 12.1% of all mention traffic flows to broadcast/parasocial accounts versus 1.0% to reciprocating hubs (Figure 6, right).

The architectural implication joins §5: simulated populations need a **broadcast tier** — and it is nearly free, since parasocial edges are one-directional by construction. The celebrity side requires no agent at all, only a content feed.




Figure 4: The scale ladder. Left: tier populations (log scale). Right: relative cost of one simulation tick, all-LLM versus the ladder.




Figure 5: Predicting second-half sentiment from first-half state. Simple persistence outperforms DeGroot diffusion over the real mention graph.

8 A Behavioral-Adherence Benchmark on Sentiment140

MatrAIx’s headline metric — do persona agents behave like their personas? — currently lacks a public, fixed-protocol instantiation that anyone can run. We contribute one on an existing public dataset.

Protocol (fixed). Cohort: the 6,245 Sentiment140 users with ≥ 20 tweets. Per user, order tweets by time; train = first 80%, test = last 20% (minimum 4). Metric: fraction of the 42,998 held-out tweets whose sentiment label the persona model predicts, reported overall and as the per-user distribution. Rules: no access to any user’s test tweets at training or persona-construction time; train-tweet text is unrestricted (features, pretrained classifiers, persona-conditioned LLM prompting, fine-tuning).

Baseline. The zero-parameter persona — predict each user’s train-majority sentiment — attains **63.1%** (chance: 50%, label-balanced corpus). Figure 8 shows the per-user distribution: bimodal, with a large mass of highly consistent users and a long tail of volatile

ones for whom persona-only prediction fails. This is the floor a surrogate tier starts from and an LLM tier must climb.

The bar. MatrAIx reports 91.5% behavioral adherence across ten attributes in four environments. The protocols are not comparable — theirs measures declared-attribute expression in synthetic interactions, ours measures real held-out behavior — so we frame 91.5% as an aspirational bar on our protocol, not a record to break on equal terms. The open question the benchmark poses: how far above 63.1% can persona-conditioned models go on real, noisy, distantly-labeled human behavior?

9 Limitations

Our substrate is 2009 Twitter: 140-character posts, no quote tweets, no algorithmic feed; behavioral constants observed here need not transfer to modern platforms. Sentiment labels are distant (emoticon-derived) and noisy. The corpus is a partial sample, which censors reciprocity; we mitigate by restricting reciprocity claims to in-sample authors with audience ≥ 20 and labeling out-of-sample targets as unobserved. The scale-ladder cost model uses nominal relative unit costs, not measured wall-clock inference costs. The diffusion experiment tests one model (DeGroot, $\alpha = 0.7$, 10 steps) at one horizon; richer contagion models could yet beat persistence. Finally, the adherence benchmark measures a single behavioral dimension (sentiment); MatrAIx-style multi-attribute adherence on this cohort is natural future work.

10 Conclusion

The methodology of billion-agent social simulators is not confined to billion-agent budgets. On one machine




Figure 6: Left: relationship classes among the 293,212 pairs with observable reciprocity. Right: share of total mention volume by target class; parasocial + broadcast attention outweighs reciprocating community hubs by an order of magnitude.




Figure 7: The benchmark section of the interactive explainer (<https://popsim.micahstubbs.ai/#adherence>): the 63.1% zero-parameter baseline against the 50% chance marker, and the per-persona accuracy distribution over held-out tweets.

and one public dataset we reproduced the persona-grounding and mixture-of-models patterns of Light Society and MatraIx, validated a diffusion assumption against ground truth (and found it wanting), mapped the parasocial structure a realistic simulator must include, and left behind a fixed, public behavioral-adherence benchmark with a 63.1% baseline and a 91.5% bar. The gap between those numbers is, we hope, an invitation.

Artifact Availability

The interactive explainer shown in Figures 1 and 7 is live at <https://popsim.micahstubbs.ai>. All analysis code, data pipelines, the benchmark protocol, and the source of this paper are at <https://github.com/micahstubbs/pop-world-model-sim>; the deployed site source is at <https://github.com/micahstubbs/>

popsim.

Acknowledgments

Built at the Postlabor.dev “Building Economic World Models” Simulation Build Night (August 2026). We thank Jake Schwartz for collaboration during the build night, and the event hosts Clovis Vinant-Tang (<https://x.com/ClovisVinant>), Tanzeela Khan (<https://x.com/tanzealous>), and Yvonne Chen (https://x.com/ycyvonne_) for organizing the evening that made this work possible.

References

- [1] H. Guan, J. He, L. Fan, Z. Ren, S. He, X. Yu, Y. Chen, X. Xu, S. Zheng, Y. Gao, E. Chen, T.-Y. Liu, Z. Liu.




Figure 8: Per-user adherence of the zero-parameter persona baseline on held-out tweets. Overall: 63.1% over 42,998 tweets.

- Modeling Earth-Scale Human-Like Societies with One Billion Agents. *arXiv:2506.12078*, 2025.
- [2] X. Li, Y. Hao, J. Hou, J. Huang, et al. MatrAIx: Simulating the World with 8.3 Billion Persona Agents. *arXiv:2608.04205*, 2026.
 - [3] J. S. Park, J. O’Brien, C. J. Cai, M. R. Morris, P. Liang, M. S. Bernstein. Generative Agents: Interactive Simulacra of Human Behavior. In *Proc. UIST*, 2023. *arXiv:2304.03442*.
 - [4] A. Go, R. Bhayani, L. Huang. Twitter Sentiment Classification using Distant Supervision. Stanford CS224N Project Report, 2009.
 - [5] F. Barbieri, J. Camacho-Collados, L. Espinosa-Anke, L. Neves. TweetEval: Unified Benchmark and Comparative Evaluation for Tweet Classification. In *Findings of EMNLP*, 2020. *arXiv:2010.12421*.
 - [6] H. Kwak, C. Lee, H. Park, S. Moon. What is Twitter, a Social Network or a News Media? In *Proc. WWW*, 2010.
 - [7] D. Horton, R. R. Wohl. Mass Communication and Para-Social Interaction: Observations on Intimacy at a Distance. *Psychiatry*, 19(3):215–229, 1956.
 - [8] M. H. DeGroot. Reaching a Consensus. *Journal of the American Statistical Association*, 69(345):118–121, 1974.